

In order to replicate the results reported, make sure all dependencies are installed first using pip. Results might be slightly different on different runs, or if you're running on different hardware, but it should be very similar.

Make sure "use_all_projects" variable in train.py is set to False so you can test individual projects, then set the "project" variable to whichever you want to test. Then, run train.py and wait for it to finish, results will be printed at the end of execution. On my setup this took around 5-10 minutes but may take longer on less powerful hardware.

Baseline results can be replicated by running the "br_classification.py" file in the "baseline_model" folder – this is just the lab1 solution so just change the "project" variable as necessary.