The code was executed on PyCharm 2024 IDE, running on Windows 11 Operating System.

NVIDIA GPU also was used. If you do not have access to an NVIDIA GPU, you might need to use the CPU version of PyTorch.

Ensure that you are using **Python 3.10**

To replicate the results reported, make sure all dependencies as specified in requirements.pdf are installed first using pip.

Results may be slightly different on different runs of the algorithm, or if different hardware is used.

Make sure that “use_all_projects” at the top of train.py is set to False so that you can test individual projects, then set the “project” variable to whichever dataset you want to get results for. Then, run train.py and wait for it to finish, results will be printed out to the console at the end of execution. On my setup, this process took around 5-10 minutes but it may take longer on less powerful hardware.

Baseline results can be replicated by running the “br_classification.py” file in the “baseline_model” folder – this is just the lab1 solution so just change “project” variable as necessary.